

# Xv6 Virtual Memory and Sharing

Instructor: Yaoqing Liu


TA: Garegin Grigoryan

# Shared Memory

P1 virtual memory

Physical memory

P2 virtual memory



# Download New Xv6

- ssh to odin and change directory to Lab7

```
$ssh YourName@odin.cslabs.clarkson.edu
```

```
$cd ~/cs444-s18/Lab7
```

- Download xv6.tar.gz file to your work directory  
Lab7

```
$wget http://people.clarkson.edu/~liu/CS444/  
Spring18/xv6.tar.gz
```

- Unzip it

```
$tar -xzf xv6.tar.gz
```

# What You Do

- Implement a system call called `shm_yourname`
- Allocates a shared memory segment

## SYNOPSIS

`Void *shm_yourname(int key, int num_pages)`

- if processes call **`shm_yourname`** with the same **key** for the first argument, then they will share the specified number of physical pages
- Using different keys in different calls to **`shm_yourname()`** corresponds to different physical pages

# More Description

- **shm\_yourname** returns the virtual address of the shared pages to the caller, so the process can read/write them
- **shm\_yourname** should map the shared physical pages to the next available virtual pages, starting at the high end of that process' address space
- 0x80000000 and above is reserved for the kernel

# Example

- When a process calls **shm\_yourname(0, 1)**:
  - The OS should map 1 physical page into the virtual address space of the caller, starting at the very high end of the address space
- If another process then calls **shm\_yourname(0, ANY\_VALUE)**:
  - this process should also get that same 1 page mapped into its virtual address space (possibly at a different virtual address)
  - Second argument is ignored if the key has been already existing
- The two processes can then each read and write to this page and thus communicate
- A third or fourth process calls **shm\_yourname(0, ANY\_VALUE)**
  - Same page mapped into their address spaces as well

# Another example

- Another key is used, **shm\_yourname(1, 3)**
  - This corresponds to a new shared region
- if any other process calls **shm\_yourname(1, 3)**
  - OS will map 3 (new) physical pages into the address space of the calling process
  - Associate these 3 pages with key value 1
  - Subsequent calls that use **key=1** will map these three pages into the calling process' address space
- To reduce the complexity, start num=1

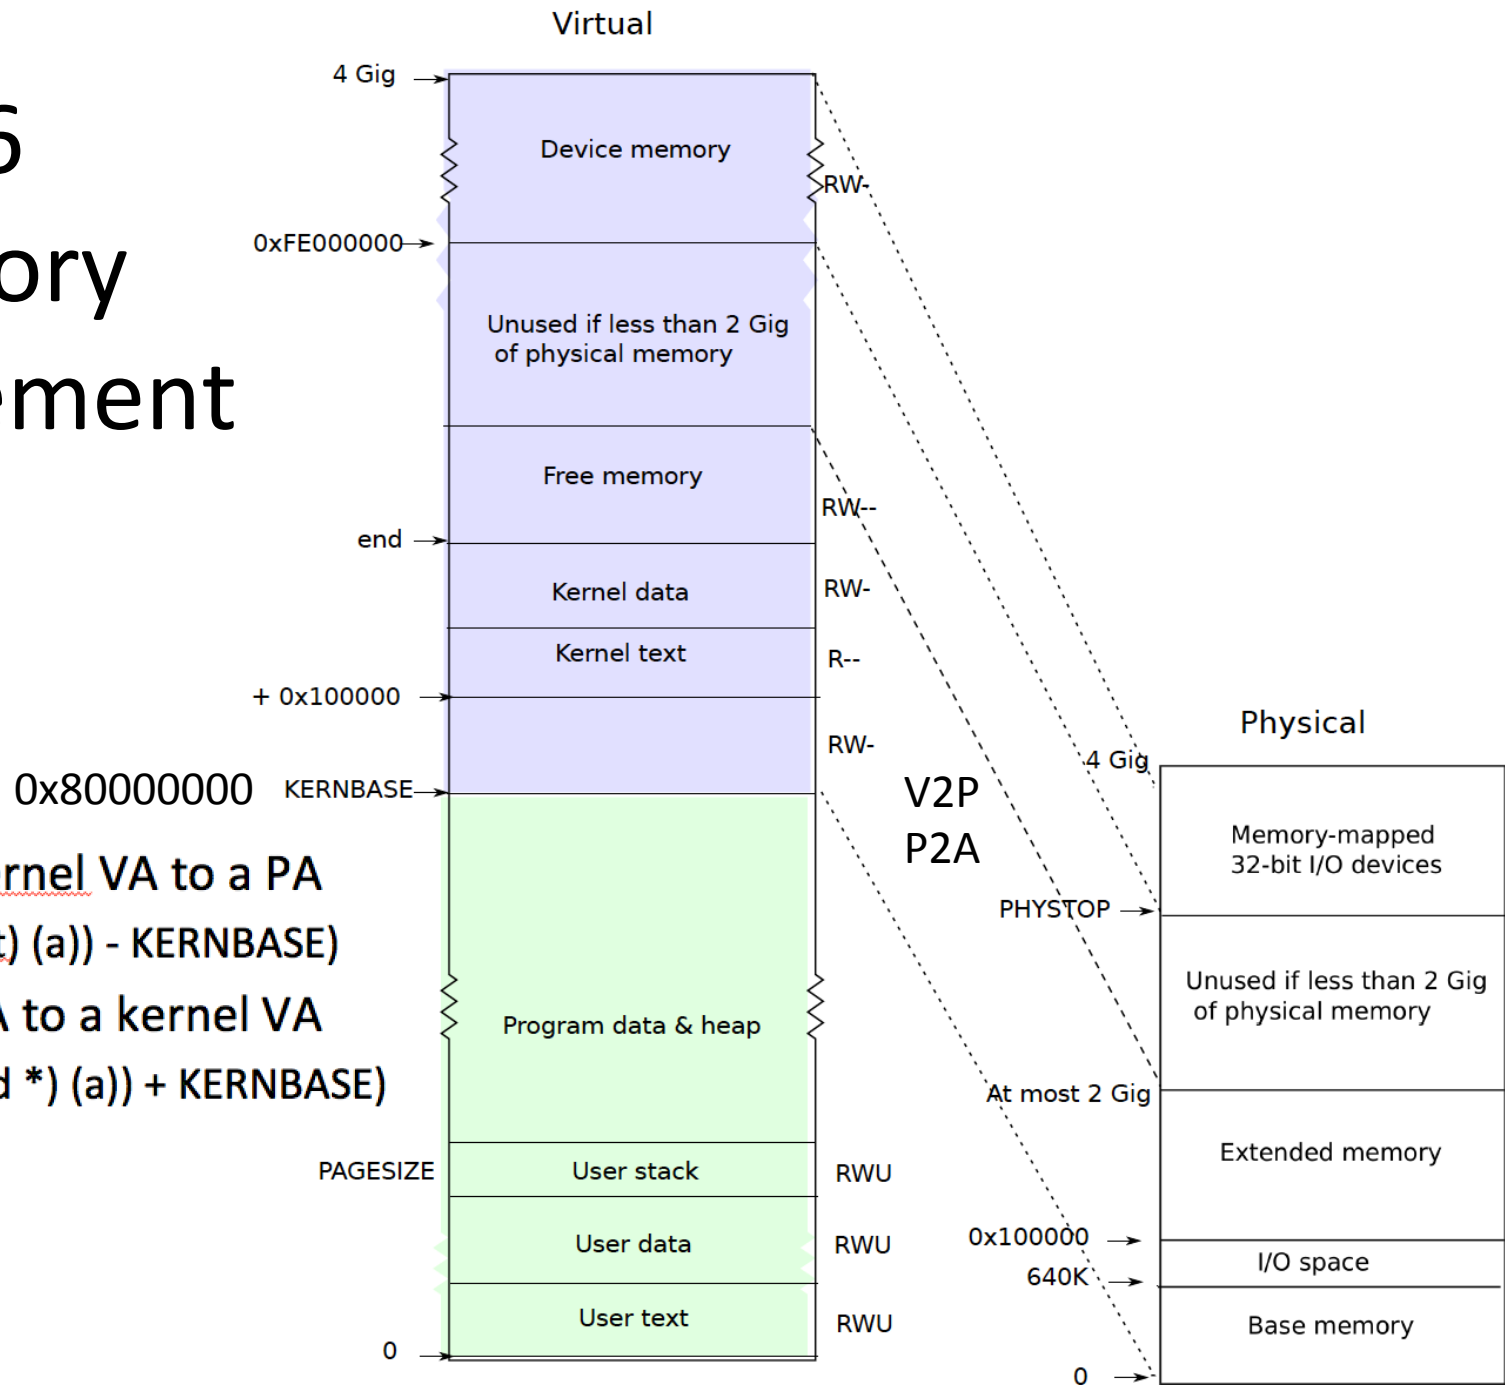
# Required Features

- Same process can call `shm_yourname` more than once
- Keys are globally visible, not for a particular process
- When a fork is called, every shared region has to be accessible to the new process **without** calling `shm_yourname()` again
- When a process exits, make sure that shared pages stay

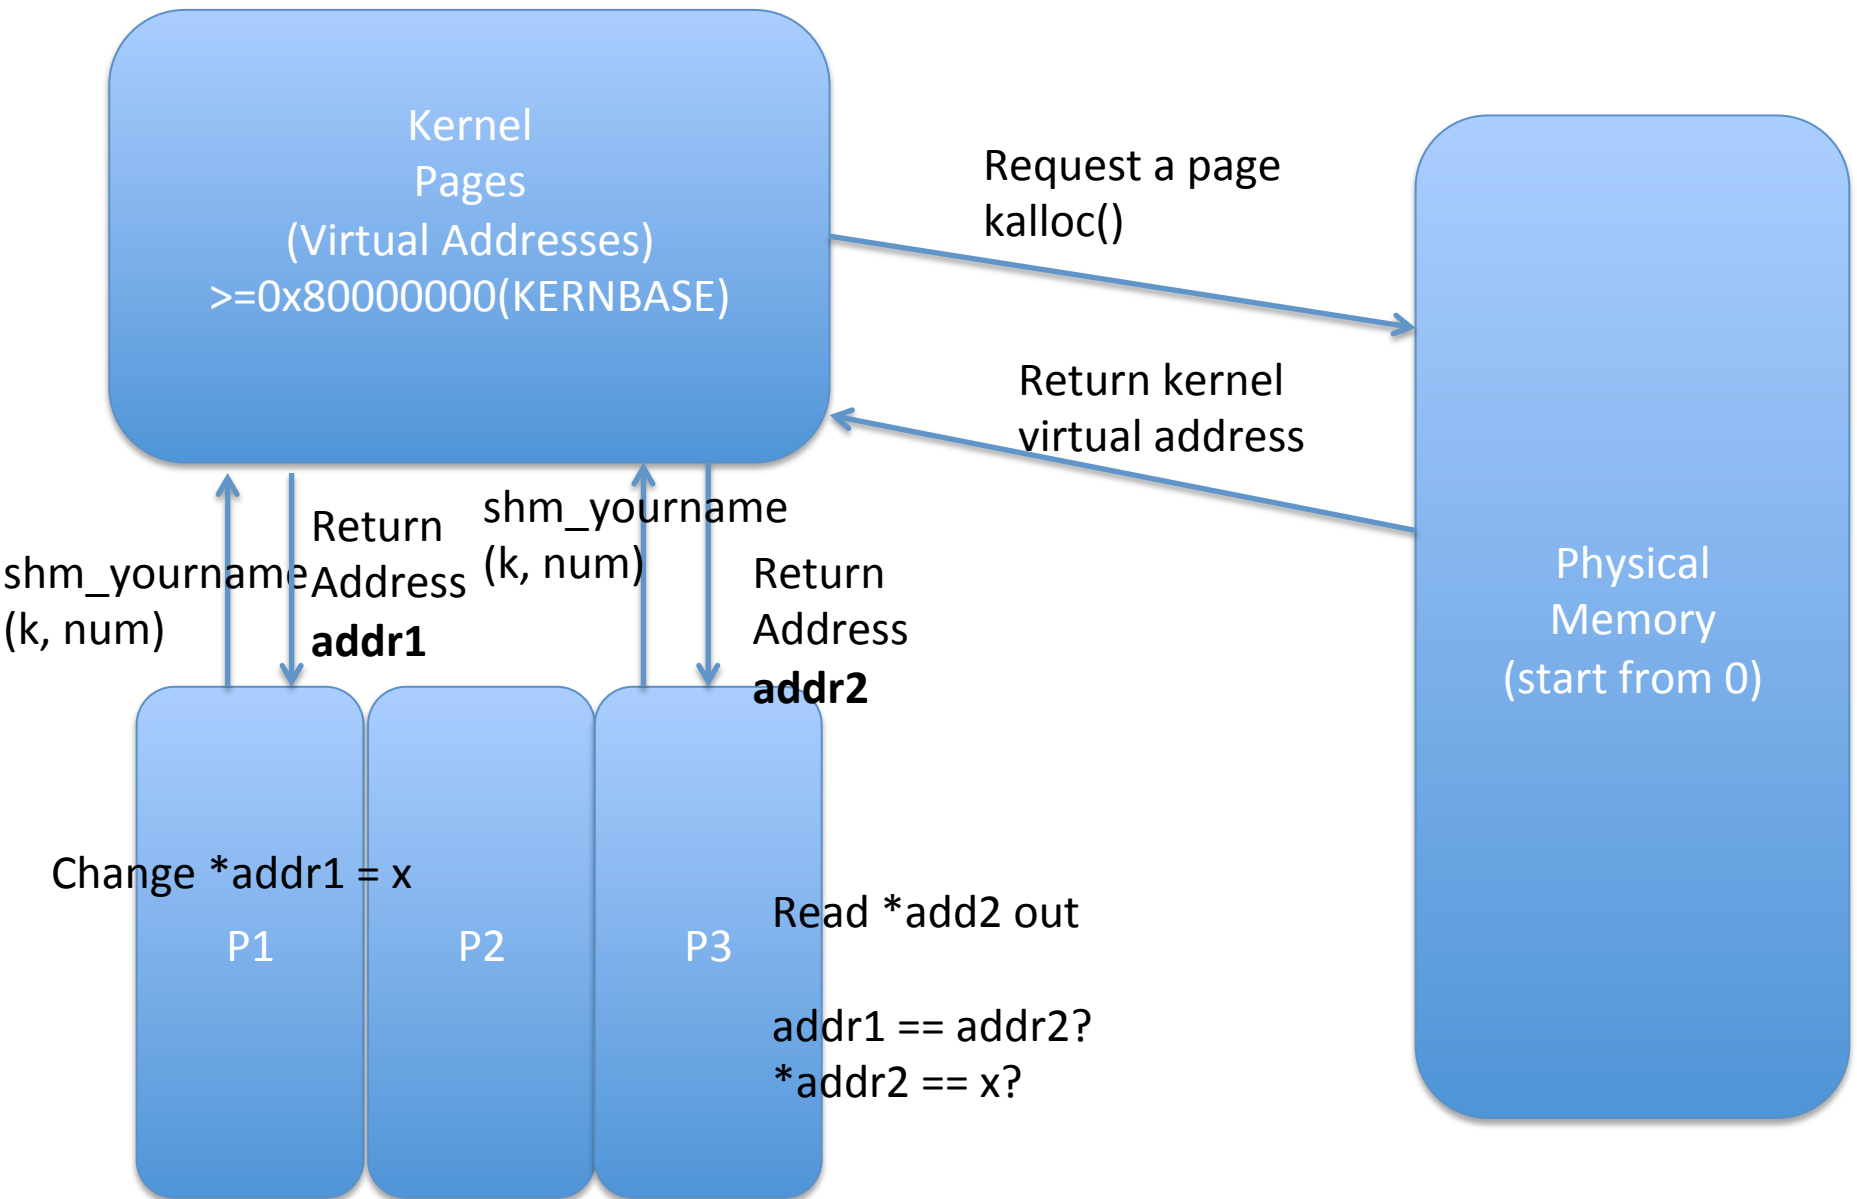


# Xv6 memory Management

- Translate a kernel VA to a PA  
–  $V2P(a) (((uint) (a)) - KERNBASE)$
- Translate a PA to a kernel VA  
–  $P2V(a) (((void *) (a)) + KERNBASE)$



# One way to do page sharing?



# One Way for Sharing Pages

## Initialization

- Define a global variable SharedMem[PAGES]
- Allocate pages (*kalloc*) using a function ShareMemInit() in vm.c
- Assign the kernel virtual addresses to SharedMem[PAGES] during booting
- Call ShareMemInit() during booting in **main.c**

## Add System Call shm\_yourname(key, num)

- Change these files: defs.h, syscall.c, syscall.h, sysproc.c, user.h
- Implement it in vm.c

## Implementation

- Start from top of virtual address in user space (0x80000000)
- The first virtual address should be the starting point of the page:  
 $0x80000000 - \text{PAGESIZE}(4K, 0x1000) = 0x7ffff000$
- Walk through the page table and map virtual addresses to their physical addresses (walkpgdir and mappages functions)
- Physical addresses can be obtained from SharedMem[PAGES] (V2P function)
- **You can use your own way to implement it**

# Testing

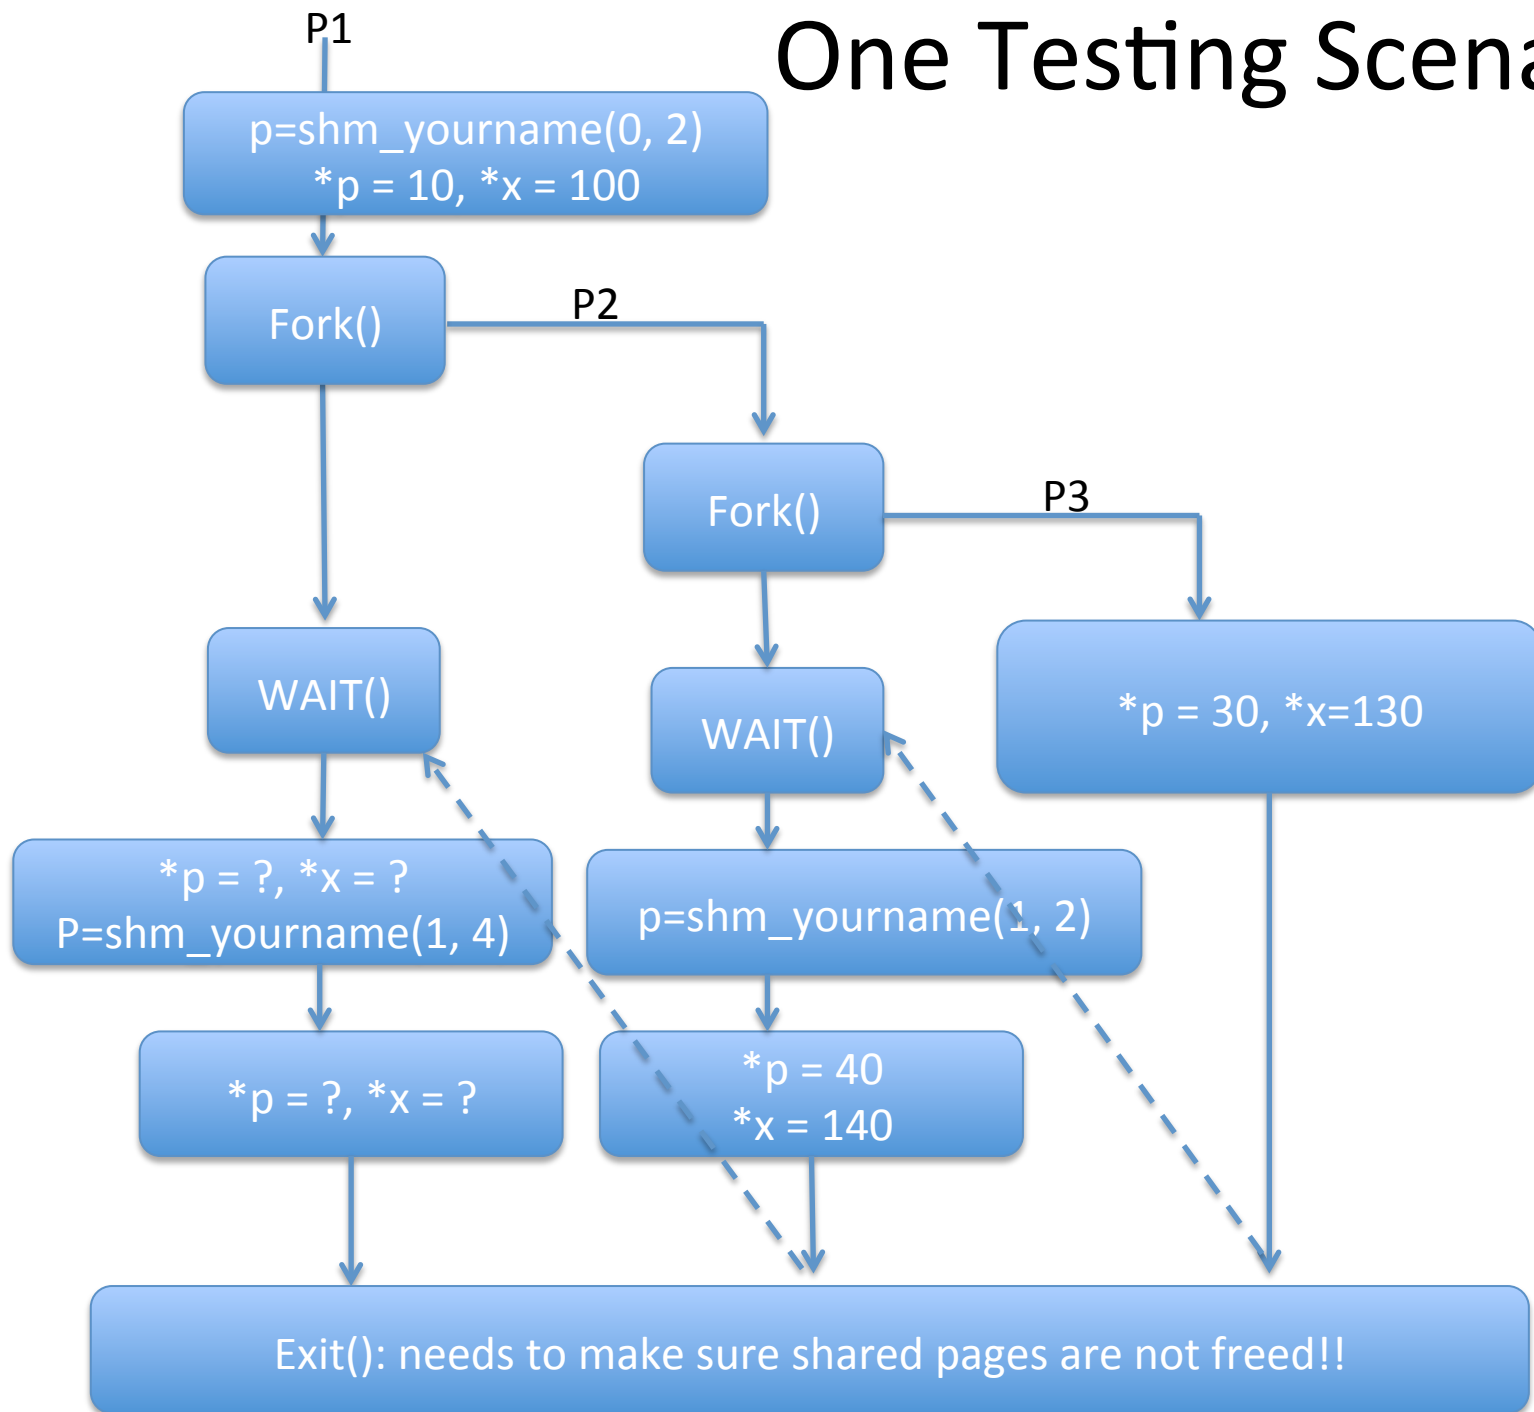
## Test Page Sharing

- Write a user program `test_shm_yourname.c` in directory `userspace` to call the system call
- `Fork()` is needed to test the shared memory between processes
- One process assigns a value to a location in the shared memory and another process read it out, the two values should be the same

## Test Key Identification

- Process 1 uses key 0
- Process 2 uses key 1 and writes to the shared page X
- process 3 uses key 1 and read the shared page value, which should be X as well

# One Testing Scenario



```

16  int p = shmgetat(0, 2);
17  int* test;
18  test = (int *)p;
19  *test = 10;
20  int t = 100;
21  int* x = &t;
22  if(fork()==0){
23      if(fork() == 0){
24          //Current limitation: key needs to be used continuously, 0 first, then 1, 2...
25          //p = shmgetat(0, 3);
26          test = (int *)p;
27          *test = 30;
28          *x = 130;
29          printf(1, "Grand Child: %x, %d, %x, %d\n", test, *test, x, *x);
30      }
31      else{
32          wait();
33          p = shmgetat(1, 2);
34          printf(1, "Child sharing key 1\n");
35          test = (int *)p;
36          *test = 40;
37          *x = 140;
38          printf(1, "Child: %x, %d, %x, %d\n", test, *test, x, *x);
39      }
40  }
41  else{
42      wait();
43      printf(1, "Before sharing key 1\n");
44      printf(1, "Parent: %x, %d, %x, %d\n", test, *test, x, *x);
45      p = shmgetat(1, 4);
46      printf(1, "After sharing key 1\n");
47      test = (int *)p;
48      printf(1, "Parent: %x, %d, %x, %d\n", test, *test, x, *x);
49  }
50  exit();

```

# Results

```
$ shmgetat
```

```
This page will be shared:0, 8dfbb000, 7ffff000
```

```
This page will be shared:1, 8dfba000, 7fffe000
```

```
Grand Child: 7FFFE000, 30, 2FBC, 130
```

```
This page will be shared:2, 8dfb9000, 7fffd000
```

```
This page will be shared:3, 8dfb8000, 7fffc000
```

```
Child sharing key 1
```

```
Child: 7FFFC000, 40, 2FBC, 140
```

```
Before sharing key 1
```

```
Parent: 7FFFE000, 30, 2FBC, 100
```

```
This page will be shared:2, 8dfb9000, 7fffd000
```

```
This page will be shared:3, 8dfb8000, 7fffc000
```

```
After sharing key 1
```

```
Parent: 7FFFC000, 40, 2FBC, 100
```

# Usage of Functions

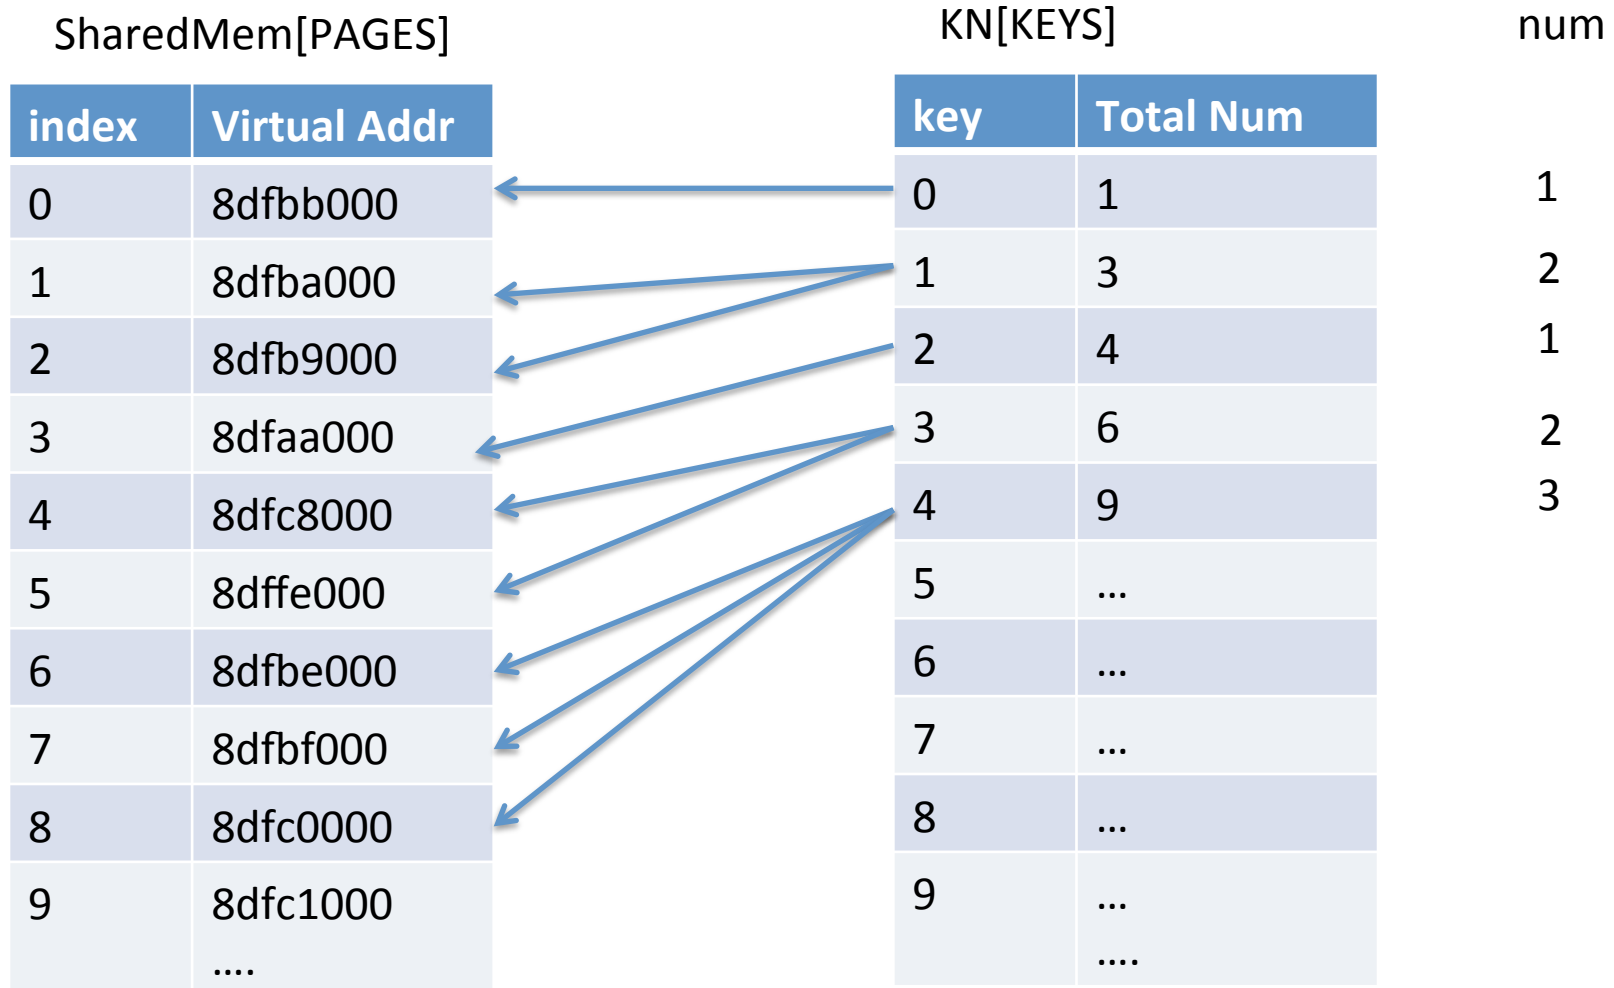
- `V2P(a) (((uint) (a)) – KERNBASE)`
  - Translate a kernel VA to a PA
- `P2V(a) (((void *) (a)) + KERNBASE)`
  - Translate a PA to a kernel VA
- `PTE_ADDR(pte *pte)`
  - Translate from a PTE to physical page number
- `Char* kalloc(void)`
  - Allocate one 4096-byte page of physical memory
  - Returns a pointer that the kernel can use
  - Returns 0 if the memory cannot be allocated
- `static int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)`
  - Create PTEs for virtual addresses starting at va that refer to physical addresses starting at pa
  - va and size might not be page-aligned
  - Perm: PTE\_U, PTE\_W, PTE\_P...



# More Functions

- `walkpgdir(pde_t *pgdir, const void *va, int alloc)`
  - Search a PTE or create a new PTE
  - If `allocate == 0`, return the address of the PTE in page table `pgdir` that corresponds to virtual address `va`
  - If `alloc!=0`, create a new page and initialize a new PTE based on `va`, set the new PTE with permission `PTE_P`, `PTE_W` and `PTE_U`
- `Int deallocvm(pde_t *pgdir, uint oldsz, uint newsz)`
  - Deallocate user pages to bring the process size from `oldsz` to `newsz`
  - **You need to modify this function to make sure shared pages won't be freed when one process exits**
- `pde_t* copyvm(pde_t *pgdir, uint sz)`
  - Given a parent process's page table, create a copy of it for a child
  - Modify it to make sure a child can inherit shared pages from parent

# One way to implement key



# Work items to enable keys

- Define a global variable `g_k` and a global array `KN[KEYS]` to track which pages have been shared
- Manipulate `g_k` and `KN` to map their items to allocated memory pages `ShareMem[PAGES]`
- Key Num `KN[key]`

0	3	3
1	2	5
2	1	6
3	1	7

`KN[key]` keeps the total number of pages up to now shared by different processes, the values can be used as indices to track which items in `ShareMem[PAGES]` have been shared.

For example, key 3 has one item shared, which is corresponding to index 6 in `ShareMem[PAGES]`

# Submission

- Capture screenshots for source code, compiling process, and results
- Combine them into a pdf file and submit it to moodle
- Due: March 26 (Monday), 11:55pm